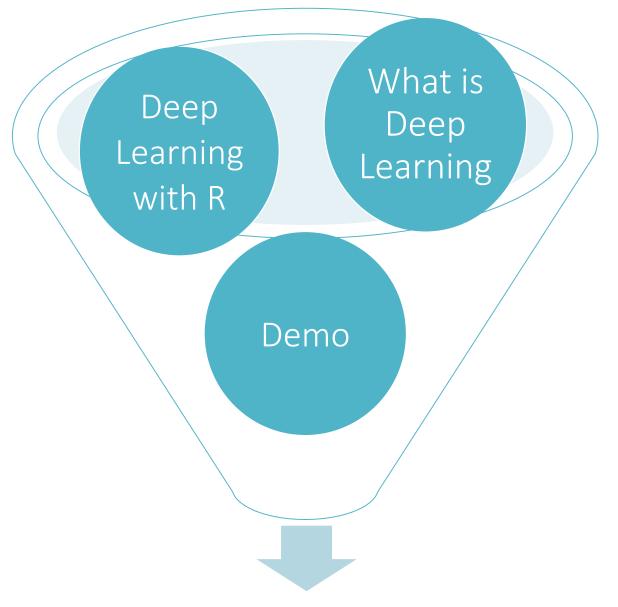
Deep Learning with R

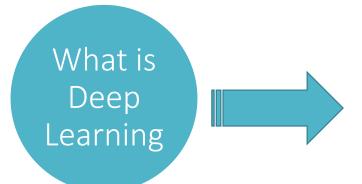
Francesca Lazzeri - @frlazzeri Data Scientist II - Microsoft, Al Research



Agenda



Better understanding of R DL tools

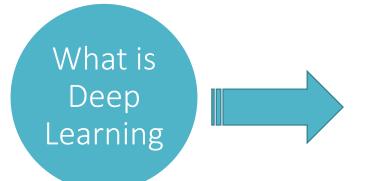


Fundamental concepts in Deep Learning

Forward
Propagation
Algorithm

Gradient Descent

Backpropagation



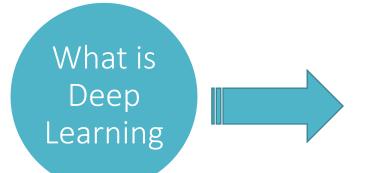
Fundamental concepts in Deep Learning

Forward Propagation Algorithm

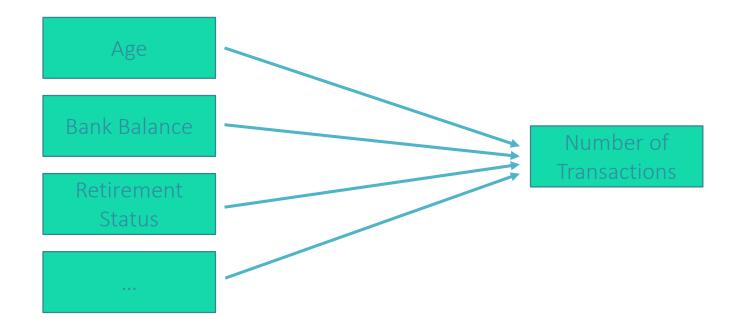
Activation Functions

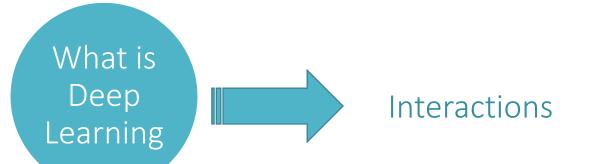
Gradient Descent

Backpropagation



Example as seen by linear regression

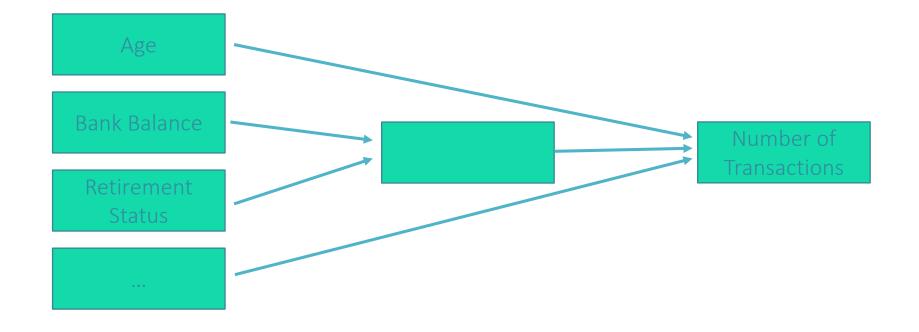




- Neural networks account for interactions really well
- Deep learning uses especially powerful neural networks for:
 - * Text
 - * Images
 - * Videos
 - * Audio
 - * Source code

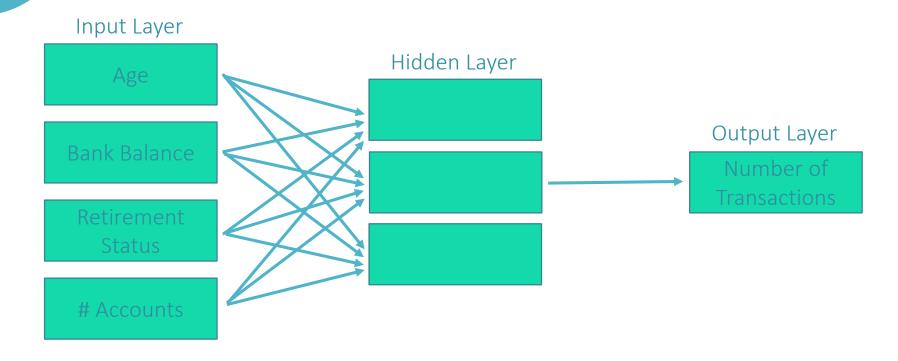


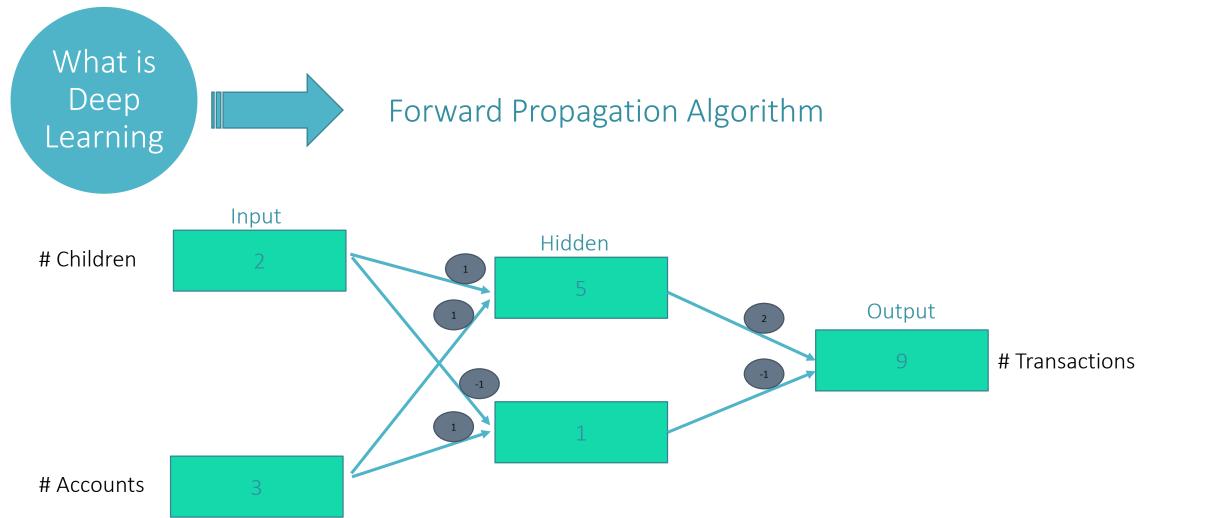
Deep learning models capture interactions

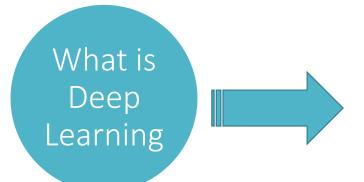


What is
Deep
Learning

Interactions in neural networks







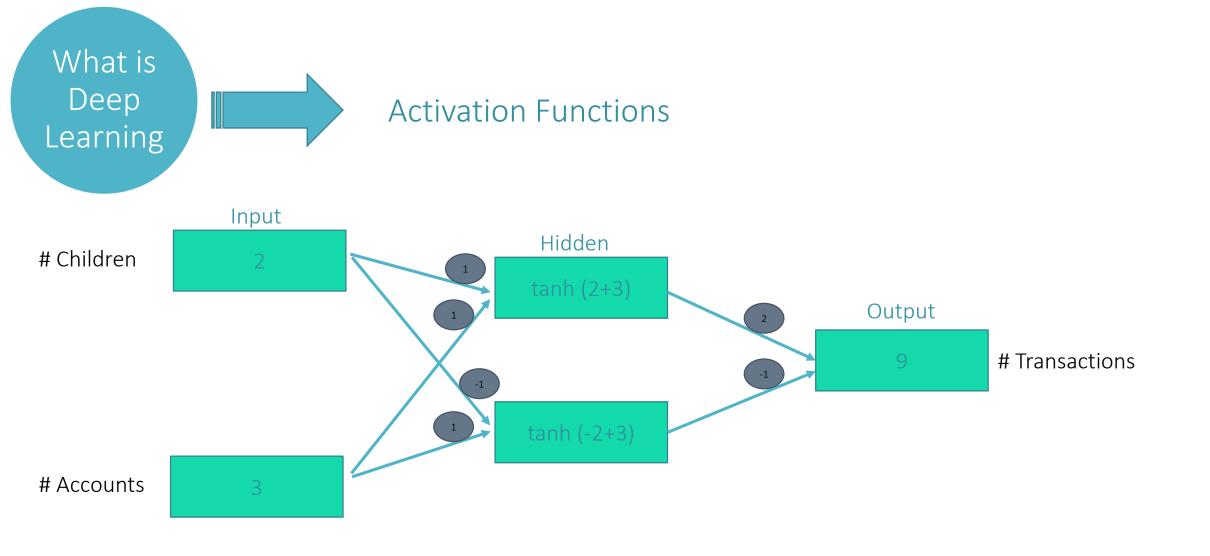
Fundamental concepts in Deep Learning

Forward Propagation Algorithm

Activation Functions

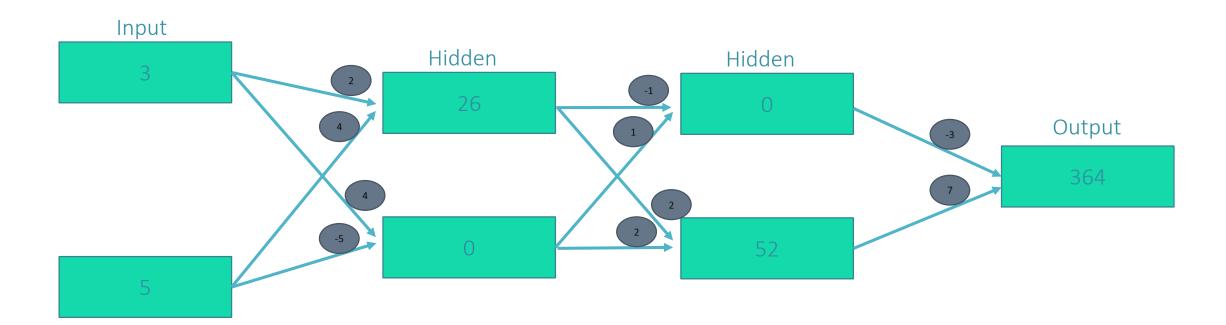
Gradient Descent

Backpropagation





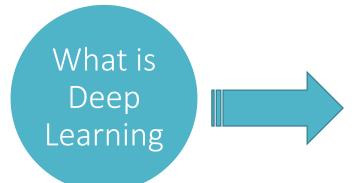
ReLU Activation Function





Representation Learning

- Deep networks internally build representations of patterns in the data
- Partially replace the need for feature engineering
- Subsequent layers build increasingly sophisticated representations of raw data
- Modeler doesn't need to specify the interactions
- When you train the model, the neural network gets weights that find the relevant patterns to make better predictions



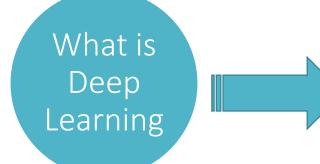
Fundamental concepts in Deep Learning

Forward Propagation Algorithm

Activation Functions

Gradient Descent

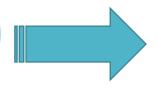
Backpropagation



The Need for Optimization

- Predictions with multiple points
 - * Making accurate predictions gets harder with more points
 - * At any set of weights, there are many values of the error
 - * Correspond to the many points we make predictions for
- Loss function
 - * Aggregate errors in predictions from many data points into single number
 - * Measure of model's predictive performance

What is Deep Learning



The Need for Optimization

Squared error loss function

Prediction	Actual	Error	Squared Error
10	20	-10	100
8	3	5	25
6	1	5	25

o Total Squared Error: 150

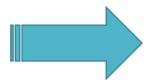
Mean Squared Error: 50

Lower loss function value means a better model

o Goal: find the weights that give the lowest value for the loss function

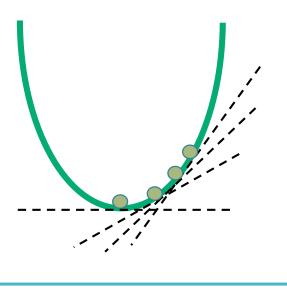
o Gradient descent!

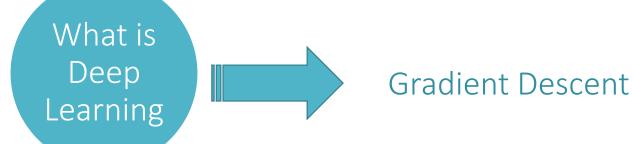
What is
Deep
Learning



Gradient Descent

Loss(w)





Slope calculation example



- To calculate the slope for a weight, need to multiply:
 - * Slope of the loss function w.r.t value at the node we feed into
 - * The value of the node that feeds into our weight
 - * Slope of activation function w.r.t value we feed into



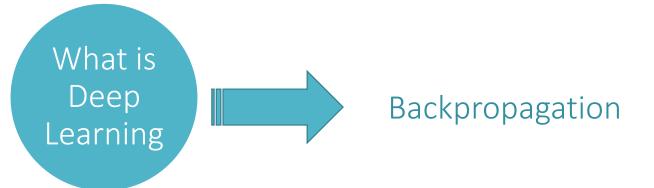
Fundamental concepts in Deep Learning

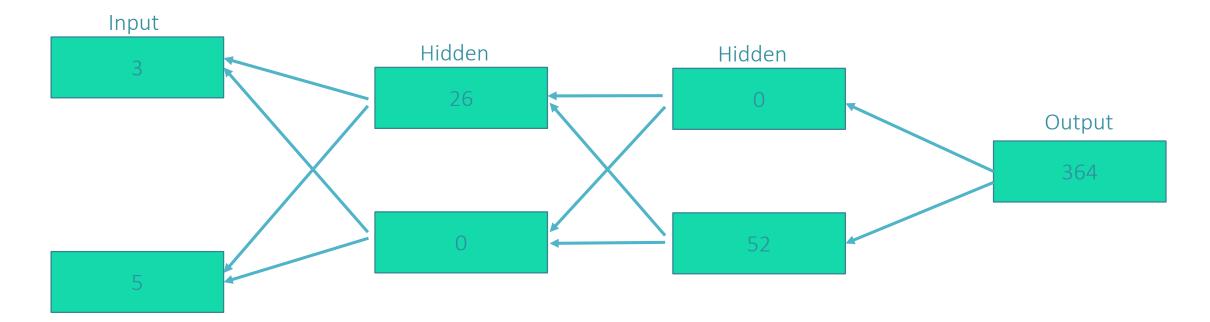
Forward Propagation Algorithm

Activation Functions

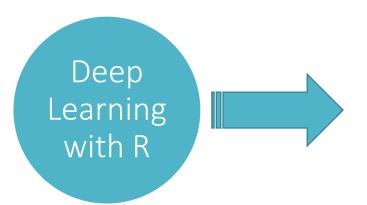
Gradient Descent

Backpropagation





- Allows gradient descent to update all weights in neural network (by getting gradients for all weights)
- o Go back one layer at a time
- o Important to understand the process, but you will generally use a library that implements this



MXNetR

- Feed-forward neural network
- Convolutional neural network (CNN)

darch

- Restricted Boltzmann machine
- Deep belief network

deepnet

- Feed-forward neural network
- Restricted Boltzmann machine
- Deep belief network
- Stacked autoencoders

H2O

- Feed-forward neural network
- Deep autoencoders

deepr

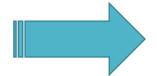
- Simplify some functions from H2O
- Deepnet packages

Deep Learning with R

Mo
МХ
MX
C
dar
de
de

	Model/Dataset	MNIST		Iris		Forest Cover Type	
		Accuracy (%)	Runtime (sec)	Accuracy (%)	Runtime (sec)	Accuracy (%)	Runtime (sec)
	MXNetR (CPU)	98.33	147.78	83.04	1.46	66.8	30.24
	MXNetR (GPU)	98.27	336.94	84.77	3.09	67.75	80.89
	darch 100	92.09	1368.31	69.12	1.71	_	_
	darch 500/300	95.88	4706.23	54.78	2.1	_	-
	deepnet DBN	97.85	6775.4	30.43	0.89	14.06	67.97
	deepnet DNN	97.05	2183.92	78.26	0.42	26.01	25.67
	H2O	98.08	543.14	89.56	0.53	67.36	5.78
	Random Forest	96.77	125.28	91.3	2.89	86.25	9.41





R interface to Keras

Keras for R

Home

Articles ▼

Examples

Reference

News

R interface to Keras

Overview

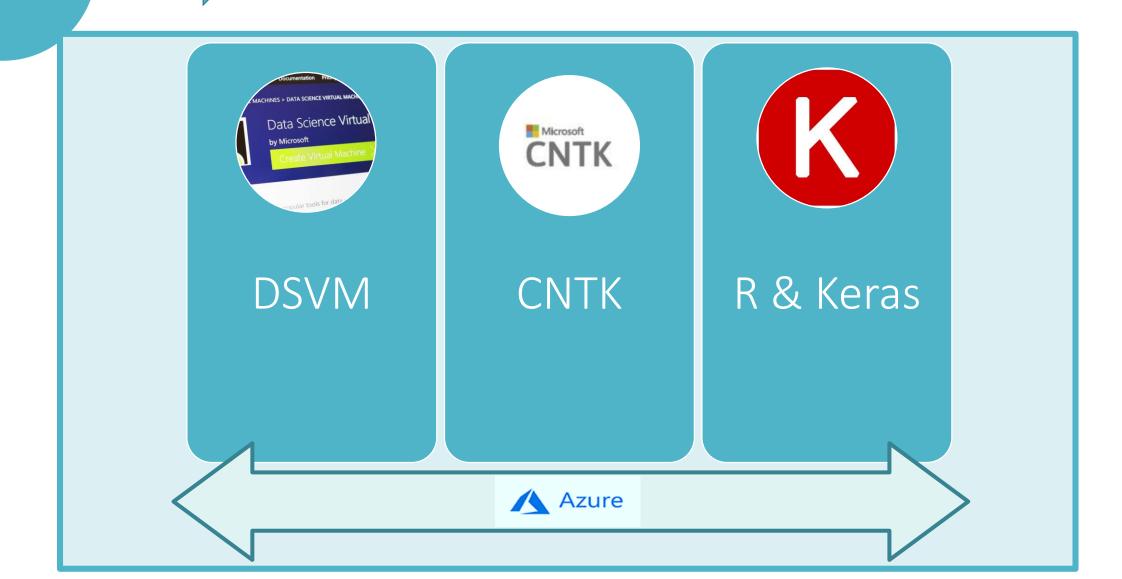
Keras is a high-level neural networks API developed with a focus on enabling fast experimentation. *Being able to go from idea to result with the least possible delay is key to doing good research.* Keras has the following key features:

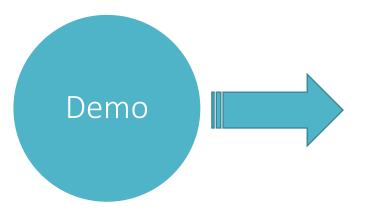
- Allows the same code to run on CPU or on GPU, seamlessly.
- User-friendly API which makes it easy to quickly prototype deep learning models.
- Built-in support for convolutional networks (for computer vision), recurrent networks (for sequence processing), and any
 combination of both.
- Supports arbitrary network architectures: multi-input or multi-output models, layer sharing, model sharing, etc. This means that Keras is appropriate for building essentially any deep learning model, from a memory network to a neural Turing machine.
- Is capable of running on top of multiple back-ends including TensorFlow, CNTK, or Theano.

This website provides documentation for the R interface to Keras. See the main Keras website at https://keras.io for additional information on the project.

Demo

Deep Learning with R on Azure with Keras and CNTK



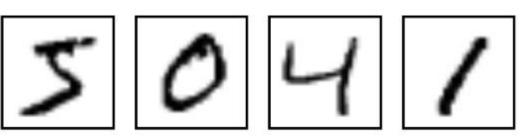


Keras Workflow Steps to Build your Model

- Specify architecture
- o Compile the model
- o Fit the model
- o Predict



Preparing the Data



```
library(keras)
install_keras()
library(keras)
mnist <- dataset_mnist()</pre>
x train <- mnist$train$x</pre>
y train <- mnist$train$y</pre>
x test <- mnist$test$x</pre>
y test <- mnist$test$y</pre>
# reshape
x_train < array_reshape(x_train, c(nrow(x_train), 784))</pre>
x_test <- array_reshape(x_test, c(nrow(x_test), 784))</pre>
# rescale
x_train <- x_train / 255
x test <- x test / 255
y_train <- to_categorical(y_train, 10)</pre>
y_test <- to_categorical(y_test, 10)</pre>
```

Demo

Defining the Model

```
model <- keras_model_sequential()
model %>%
  layer_dense(units = 256, activation = 'relu', input_shape = c(784)) %>%
  layer_dropout(rate = 0.4) %>%
  layer_dense(units = 128, activation = 'relu') %>%
  layer_dropout(rate = 0.3) %>%
  layer_dense(units = 10, activation = 'softmax')
```



Defining the Model

summary(model)

Model

Layer (type)	Output Shape	Param #
dense_1 (Dense)	(None, 256)	200960
dropout_1 (Dropout)	(None, 256)	0
dense_2 (Dense)	(None, 128)	32896
dropout_2 (Dropout)	(None, 128)	0
dense_3 (Dense)	(None, 10)	1290

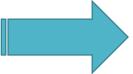
Total params: 235,146

Trainable params: 235,146 Non-trainable params: 0

Defining the Model

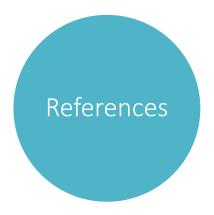
```
model %>% compile(
  loss = 'categorical_crossentropy',
  optimizer = optimizer_rmsprop(),
  metrics = c('accuracy')
)
```

Demo



Training and Evaluation

```
history <- model %>% fit(
  x_train, y_train,
  epochs = 30, batch_size = 128,
  validation split = 0.2
model %>% evaluate(x_test, y_test)
$loss
[1] 0.1149
$acc
[1] 0.9807
model %>% predict_classes(x_test)
  [1] 7 2 1 0 4 1 4 9 5 9 0 6 9 0 1 5 9 7 3 4 9 6 6 5 4 0 7 4 0 1 3 1 3 4 7 2 7 1 2
 [40] 1 1 7 4 2 3 5 1 2 4 4 6 3 5 5 6 0 4 1 9 5 7 8 9 3 7 4 6 4 3 0 7 0 2 9 1 7 3 2
 [79] 9 7 7 6 2 7 8 4 7 3 6 1 3 6 9 3 1 4 1 7 6 9
 [ reached getOption("max.print") -- omitted 9900 entries ]
```



http://blog.revolutionanalytics.com/2017/08/keras-and-cntk.html

http://www.rblog.uni-freiburg.de

https://keras.rstudio.com/

https://campus.datacamp.com/courses/deep-learning

http://gluon.mxnet.io

Thank You!

Francesca Lazzeri - @frlazzeri Data Scientist II - Microsoft, Al Research

